
Multiple Application Environments

with Red Hat Advanced Cluster Management

HighVail Systems Inc.



© 2020 version: 1.0

Contents

1	Multiple Application Environments with Red Hat Advanced Cluster Management	3
1.1	Cluster Components	4
1.2	Application Life-cycle	4
1.3	Git Repository	5
1.4	Deploying an Application to Multiple Environments	5
1.5	Review	6
1.5.1	Channel	6
1.5.2	Namespace	7
1.5.3	PlacementRule	7
1.5.4	Subscription	8
1.5.5	Application	8
2	Part 1 - Deploying an Application to Multiple Environments	9
2.1	Local Environment	9
2.2	Deploying the Application to the Development Environment	9
2.3	Deploying the application to the production environment	12
3	Part 2 - Blue/Green Deployments	15
3.1	Blue / Green Deployments on Red Hat Advanced Cluster Management	15
3.2	Local Environment	15
3.3	Upgrading the Application in the Development Environment	17
3.4	Upgrading the application on the production environment	20
4	Part 3 - Application Migration	22
4.1	Application Migration on Advanced Cluster Management	22
4.2	Local Environment	22
4.3	Creating New PlacementRules and Subscriptions	22
4.4	Migrating our Application	24
5	Part 4 - Disaster Recovery	25
5.1	Disaster Recovery on Red Hat Advanced Cluster Management	25
5.2	Local Environment	26
5.3	Configuring the Required Advanced Cluster Management Manifests	26
5.3.1	Configuring the Subscription to use the New PlacementRule	27

1 Multiple Application Environments with Red Hat Advanced Cluster Management

Based on the blog by Mario Vázquez [source](#)

Red Hat Advanced Cluster Management features complete Application Life-cycle capabilities. Below we will explore some of the GitOps capabilities built into the Red Hat Advanced Cluster Management.

The following is the environment used.

NOTE: The labels used within this example are reused throughout the demonstration.

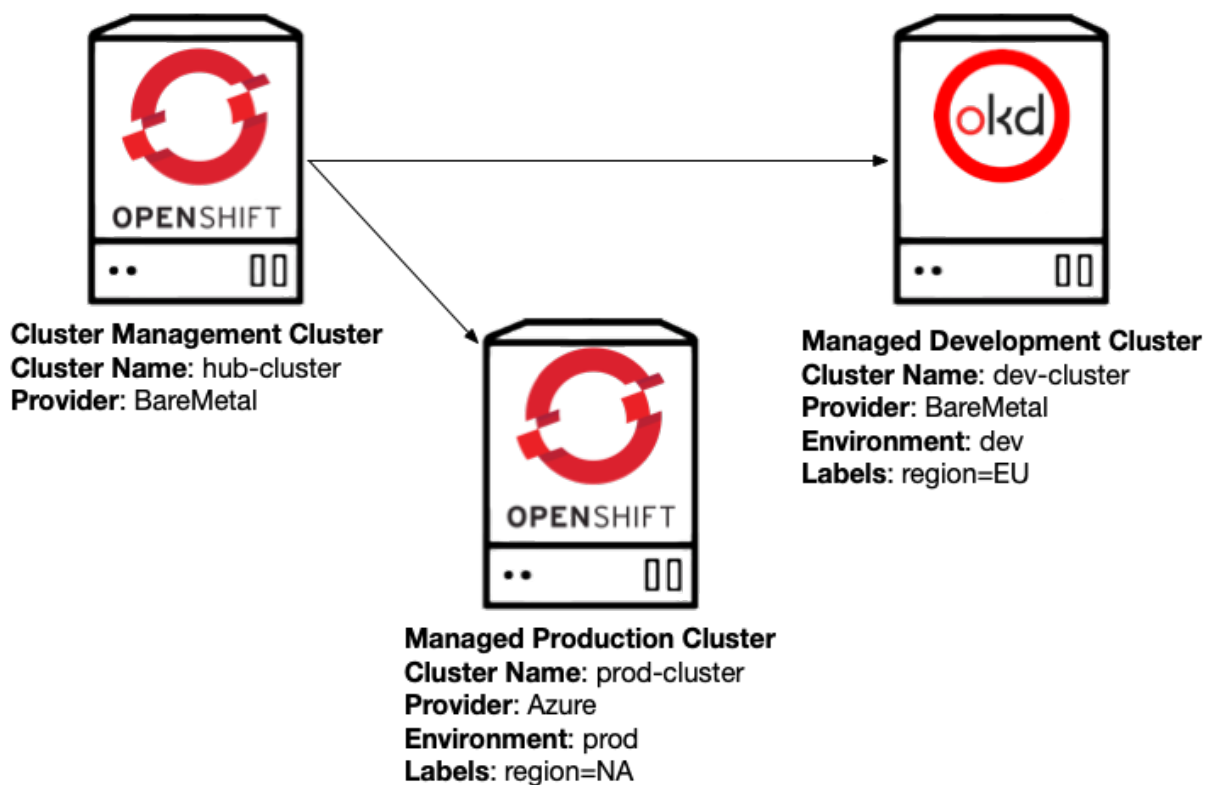


Figure 1: Cluster Architecture

Three OpenShift clusters are used for this demonstration. Red Hat Advanced Cluster Management uses a hub cluster/managed cluster model for managing the different connected clusters, the hub cluster is where Red Hat Advanced Cluster Management is running, the managed clusters are the clusters that are managed by Red Hat Advanced Cluster Management.

1.1 Cluster Components

Component	Version
Red Hat OpenShift Container Platform	4.5.11
OKD OpenShift Container Platform	4.5.0-0.okd.2020-08-12-020541
Red Hat Advanced Cluster Management for Kubernetes	2.0.3

The managed clusters have labels, which will be used later for placing applications across the different environments.

The managed development cluster is named `dev-cluster` and it is deployed in the *EU* region in a three node (all master) bare metal cluster. We also have a managed production cluster named `prod-cluster` deployed in the *NA* region on a five node (three managers and two worker nodes) Azure cluster.

1.2 Application Life-cycle

The possibilities offered by Red Hat Advanced Cluster Management with regards to application lifecycle are plentiful. In this series of blog posts we are going to focus on GitOps capabilities for the following use cases:

- Deploying an application to multiple environments
- Blue-green deployment
- Application migration
- Disaster recovery

Concepts for GitOps and Red Hat Advanced Cluster Management.

1. Channels

Channels point to a physical place where resources are stored for deployment. For this write-up we are going to use Channels of type `Git`, there are different channel types for Helm, Namespaces, etc. [Learn More](#)

2. PlacementRules

You can create and manage placement rules to define where subscriptions to Kubernetes resources and Helm releases are deployed. Use placement rules to help you facilitate multi-cluster deployments of your Kubernetes resources. [Learn More](#)

3. Subscriptions

Subscriptions serve as sets of definitions for identifying Kubernetes resources within channels by using annotations, labels, and versions. Subscription resources are defined on the hub and propagated to the managed cluster(s). The subscription controller watches the source location (channel) for new or updated resources. When a new or updated Kubernetes resource is detected, the subscription controller can download the resource directly from the source location (channel) to managed clusters without checking the Hub cluster first (because the subscription was initially propagated). The subscription can filter the Helm releases to select a specific chart version. For this case, the subscription controller checks the version parameter to identify the Helm release (chart) version to select. [Learn More](#)

4. Applications

An Application object provides a way for you to aggregate subscriptions as a group. It provides tooling and a console with a resource that allows for the aggregation and display of all the components in the Application. [Learn More](#)

1.3 Git Repository

GitOps patterns are followed to deploy the applications, the different manifests required to deploy the applications in the different environments will be stored in a Git repository, the Git structure is defined in the following table:

Branch	Description
config	Stores the base files for the applications, which apply to every environment
prod	Stores the overlay files for the applications, which apply to production environments
stage	Stores the overlay files for the applications, which apply to staging environments

NOTE: There are multiple ways to organise the Git repository, Red Hat Advanced Cluster Management allows any method.

1.4 Deploying an Application to Multiple Environments

Red Hat Advanced Cluster Management can deploy applications to multiple environments. In the following example, the application is a simple web-service that reverses words. This web-service has

two releases, the stage release which is the version the development team is testing at the moment and the production release, which is the version the customers are using.

Red Hat Advanced Cluster Management includes support for [Kustomize](#), which allows for application configuration based on the destination environments.

The same application will for both environments, but the release will be different depending on which environment the application is being deployed to.

Deployment of the applications will be through the use of the OpenShift `oc` tool and a set of `yaml` manifests with the required configurations for Red Hat Advanced Cluster Management that define a Channel, Subscription, and PlacementRule. Everything will be done via the command line, but could also be performed from the web console, as well.

There are three contexts configured, one for each environment:

Context	Description
hub	CLI Profile connected to the HUB Cluster (where ACM is deployed)
dev	CLI Profile connected to the managed development cluster (dev-cluster)
prod	CLI Profile connected to the managed production cluster (prod-cluster)

More about CLI profiles [here](#).

1.5 Review

Review of the resources used.

1.5.1 Channel

The Channel is defined as a Git type Channel that will be used by the subscriptions in order to obtain the Kubernetes resources that deploy the application.

In this case, it is configured to get the Kubernetes resources from the Git repository at `github.com/zombiefish/acm-app-lifecycle.git`.

```
apiVersion: apps.open-cluster-management.io/v1
kind: Channel
metadata:
  name: acm-app-lifecycle
```

```
namespace: open-cluster-management
spec:
  type: Git
  pathname: https://github.com/zombiefish/acm-app-lifecycle.git
```

1.5.2 Namespace

When using Subscriptions, the namespace holding the Subscription will be propagated to the destination cluster. Here, we are creating a namespace named `reverse-words-stage` that will be propagated to the development clusters by the Subscription.

```
apiVersion: v1
kind: Namespace
metadata:
  name: reverse-words-stage
```

1.5.3 PlacementRule

The Subscriptions are propagated to a list of clusters returned by a PlacementRule. this allows the selection of a cluster or group of clusters from the different environments and make that list available to the different Subscriptions.

The PlacementRule named `development-clusters` will return all clusters which are marked as *available* and that match the label `environment: dev`, in this scenario, this PlacementRule will return the managed development cluster named `dev-cluster`.

```
apiVersion: apps.open-cluster-management.io/v1
kind: PlacementRule
metadata:
  name: development-clusters
  namespace: reverse-words-stage
spec:
  clusterConditions:
  - type: "ManagedClusterConditionAvailable"
    status: "True"
  clusterSelector:
    matchExpressions: []
```

```
matchLabels:
  environment: "dev"
```

1.5.4 Subscription

The Subscription will deploy a list of Kubernetes resources (obtained from a given Channel) into a list of clusters (obtained from a given PlacementRule). You can also define where the Kubernetes resources are located within the Git repository (Channel).

This subscription uses the Channel defined earlier, and will get the Kubernetes resources from the branch stage. Within that branch, it will look for Kubernetes resources in the apps/reversewords/path.

```
apiVersion: apps.open-cluster-management.io/v1
kind: Subscription
metadata:
  name: reversewords-dev-app-subscription
  namespace: reverse-words-stage
  labels:
    app: reversewords-dev-app
  annotations:
    apps.open-cluster-management.io/git-path: apps/reversewords/
    apps.open-cluster-management.io/git-branch: stage
spec:
  channel: open-cluster-management/acm-app-lifecycle
  placement:
    placementRef:
      kind: PlacementRule
      name: development-clusters
```

1.5.5 Application

And finally, the application itself.

```
apiVersion: app.k8s.io/v1beta1
kind: Application
metadata:
  name: reversewords-dev-app
```



```
namespace: reverse-words-stage
spec:
  componentKinds:
    - group: apps.open-cluster-management.io
      kind: Subscription
  descriptor: {}
  selector:
    matchExpressions:
      - key: app
        operator: In
        values:
          - reversewords-dev-app
```

2 Part 1 - Deploying an Application to Multiple Environments

The Application creates a topology view of the applications across the clusters that are managed by Advanced Cluster Management. To do so, the Application matches one or more Subscriptions, which ultimately will create resources in the different clusters. This method allows for audability, ie; we can trace what was created by whom and where it was created.

2.1 Local Environment

For sanity, and less complexity, set a local variable for the rw git content. In this case, set the path to DEMO.

```
$ export DEMO="https://raw.githubusercontent.com/zombiefish/acm-app-lifecycle/master/acm-mar
```

2.2 Deploying the Application to the Development Environment

1. First, create the Channel definition.

```
$ oc --context hub create \
  -f $DEMO/base/00_channel.yaml

channel.apps.open-cluster-management.io/acm-app-lifecycle created
```

2. Next, create a Namespace for storing the application manifests.

```
$ oc --context hub create \  
  -f $DEMO/reversewords-stage/00_namespace.yaml  
  
namespace/reverse-words-stage created
```

3. Now create a PlacementRule that matches the managed development clusters.

```
$ oc --context hub create \  
  -f $DEMO/reversewords-stage/01_placement_rule.yaml  
  
placementrule.apps.open-cluster-management.io/development-clusters created
```

Check the PlacementRule status, note that it matches the managed development cluster named dev-cluster:

```
$ oc --context hub -n reverse-words-stage  
  get placementrule development-clusters -o yaml  
  
<OMITTED_OUTPUT>  
status:  
  decisions:  
  - clusterName: dev-cluster  
    clusterNamespace: dev-cluster
```

4. The Subscription and the Application can be created targeting the development clusters via the PlacementRule.

```
$ oc --context hub create \  
  -f $DEMO/reversewords-stage/02_subscription-dev.yaml  
  
subscription.apps.open-cluster-management.io/  
reversewords-dev-app-subscription created  
  
$ oc --context hub create \  
  -f $DEMO/reversewords-stage/03_application-dev.yaml  
  
application.app.k8s.io/reversewords-dev-app created
```

See the following Subscription status. Note that it says propagated, which means the Subscription has been sent to the destination cluster:

```
$ oc --context hub -n reverse-words-stage get
subscription reversewords-dev-app-subscription -o yaml

<OMITTED_OUTPUT>
status:
  lastUpdateTime: "2020-10-05T15:32:00Z"
  phase: Propagated
```

5. Looking at the development cluster, see that the application is up and running.

```
$ oc --context dev -n reverse-words-stage get deployments,services,pods
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/reverse-words	1/1	1	1	75s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
reverse-words	LoadBalancer	172.30.178.135	<pending>	8080:30804/TCP	75s

NAME	READY	STATUS	RESTARTS	AGE
pod/reverse-words-6dfc7b864b-bbcjc	1/1	Running	0	75s

At this point, although the application is up, there is no access the application.

6. To access the application on the network, the service needs to be exposed.

```
$ oc --context dev create \
  -f $DEMO/reversewords-stage/04_service-route.yaml

route.route.openshift.io/reverse-words created
```

Get the address of the route just created:

```
$ oc --context dev -n reverse-words-stage get route
```

NAME	HOST/PORT
reverse-words	reverse-words-reverse-words-stage.apps.okd.highvail.com

PATH	SERVICES	PORT	TERMINATION	WILDCARD
reverse-words	http		None	

If we run the same query against production cluster, we will see that there is no application running there.

```
$ oc --context prod -n reverse-words-stage get deployments,services,pods,route

No resources found in reverse-words-stage namespace.
```

Query the application to see the deployed staging release:

```
$ curl http://reverse-words-reverse-words-stage.apps.okd.highvail.com

Reverse Words Release: Stage Release v0.0.3. App version: v0.0.3
```

2.3 Deploying the application to the production environment

1. No need to create a new Channel since it will be using the same source Git repository, but a different branch.
2. Create a Namespace for storing the application manifests.

```
$ oc --context hub create \
    -f $DEMO/reversewords-prod/00_namespace.yaml

namespace/reverse-words-prod created
```

3. Now create a PlacementRule that matches the production clusters:

```
$ oc --context hub create \_
    -f $DEMO/reversewords-prod/01_placement_rule.yaml

placementrule.apps.open-cluster-management.io/production-clusters created
```

See the following PlacementRule status. Note that it matched the managed production cluster named, prod-cluster.

```
$ oc --context hub -n reverse-words-prod get
    placementrule production-clusters -o yaml

<OMITTED_OUTPUT>
status:
  decisions:
    - clusterName: prod-cluster
      clusterNamespace: prod-cluster
```

4. The Subscription and the Application can be created now targeting the production clusters via the PlacementRule.

```
$ oc --context hub create \
  -f $DEMO/reversewords-prod/02_subscription-pro.yaml

subscription.apps.open-cluster-management.io/
reversewords-pro-app-subscription created

$ oc --context hub create \
  -f $DEMO/reversewords-prod/03_application-pro.yaml

application.app.k8s.io/reversewords-pro-app created
```

See the following Subscription status. Note that it says propagated, which means the Subscription has been sent to the destination cluster:

```
$ oc --context hub -n reverse-words-prod get
subscription reversewords-pro-app-subscription -o yaml

<OMITTED_OUTPUT>
status:
  message: Active
  phase: Propagated
```

5. Finally, look at the production cluster and see the application up and running.

```
$ oc --context prod -n reverse-words-prod get deployments,services,pods
```

NAME	READY	STATUS	RESTARTS	AGE
pod/reverse-words-646456cd7c-5lss4	1/1	Running	0	102s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
reverse-words	LoadBalancer	172.30.69.224	<pending>	8080:31860/TCP	102s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/reverse-words	1/1	1	1	102s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/reverse-words-646456cd7c	1	1	1	102s

6. Query the application and see that it deployed the production release.

```
$ oc --context prod -n reverse-words-prod create \
  -f $DEMO/reversewords-prod/04_service-route.yaml

route.route.openshift.io/reverse-words created
```

Get the address of the route we just created:

```
$ oc --context prod -n reverse-words-prod get route

NAME          HOST/PORT
reverse-words  reverse-words-reverse-words-prod.apps.ocp.highvail.com
  PATH  SERVICES      PORT  TERMINATION  WILDCARD
reverse-words  http                      None
```

And query the application and see that it deployed the staging release:

```
$ curl http://reverse-words-reverse-words-prod.apps.ocp.highvail.com

Reverse Words Release: Production release v0.0.2. App version: v0.0.2
```

7. There are different versions of the application, depending on the environment we are deploying to:

```
# Query development environment
$ curl http://reverse-words-reverse-words-stage.apps.okd.highvail.com

Reverse Words Release: Stage Release v0.0.3. App version: v0.0.3

# Query production environment
$ curl http://reverse-words-reverse-words-prod.apps.ocp.highvail.com

Reverse Words Release: Production release v0.0.2. App version: v0.0.2
```

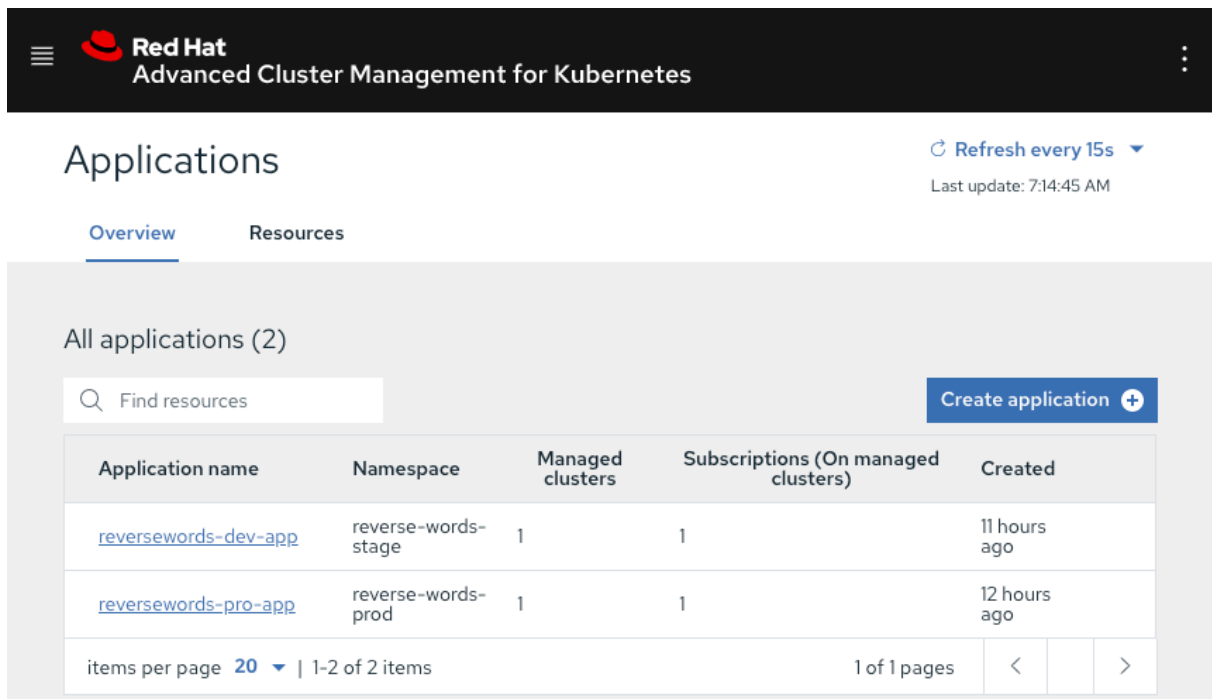
Finally, let's take a look in the web console:

ACM Applications General View

ACM Development Application View

ACM ClusterTopology View

ACM Overview



Red Hat
Advanced Cluster Management for Kubernetes

Applications

Refresh every 15s ▾
Last update: 7:14:45 AM

Overview Resources

All applications (2)

Find resources

Create application +

Application name	Namespace	Managed clusters	Subscriptions (On managed clusters)	Created
reversewords-dev-app	reverse-words-stage	1	1	11 hours ago
reversewords-pro-app	reverse-words-prod	1	1	12 hours ago

items per page 20 ▾ | 1-2 of 2 items

1 of 1 pages < >

Figure 2: ACM Application View

3 Part 2 - Blue/Green Deployments

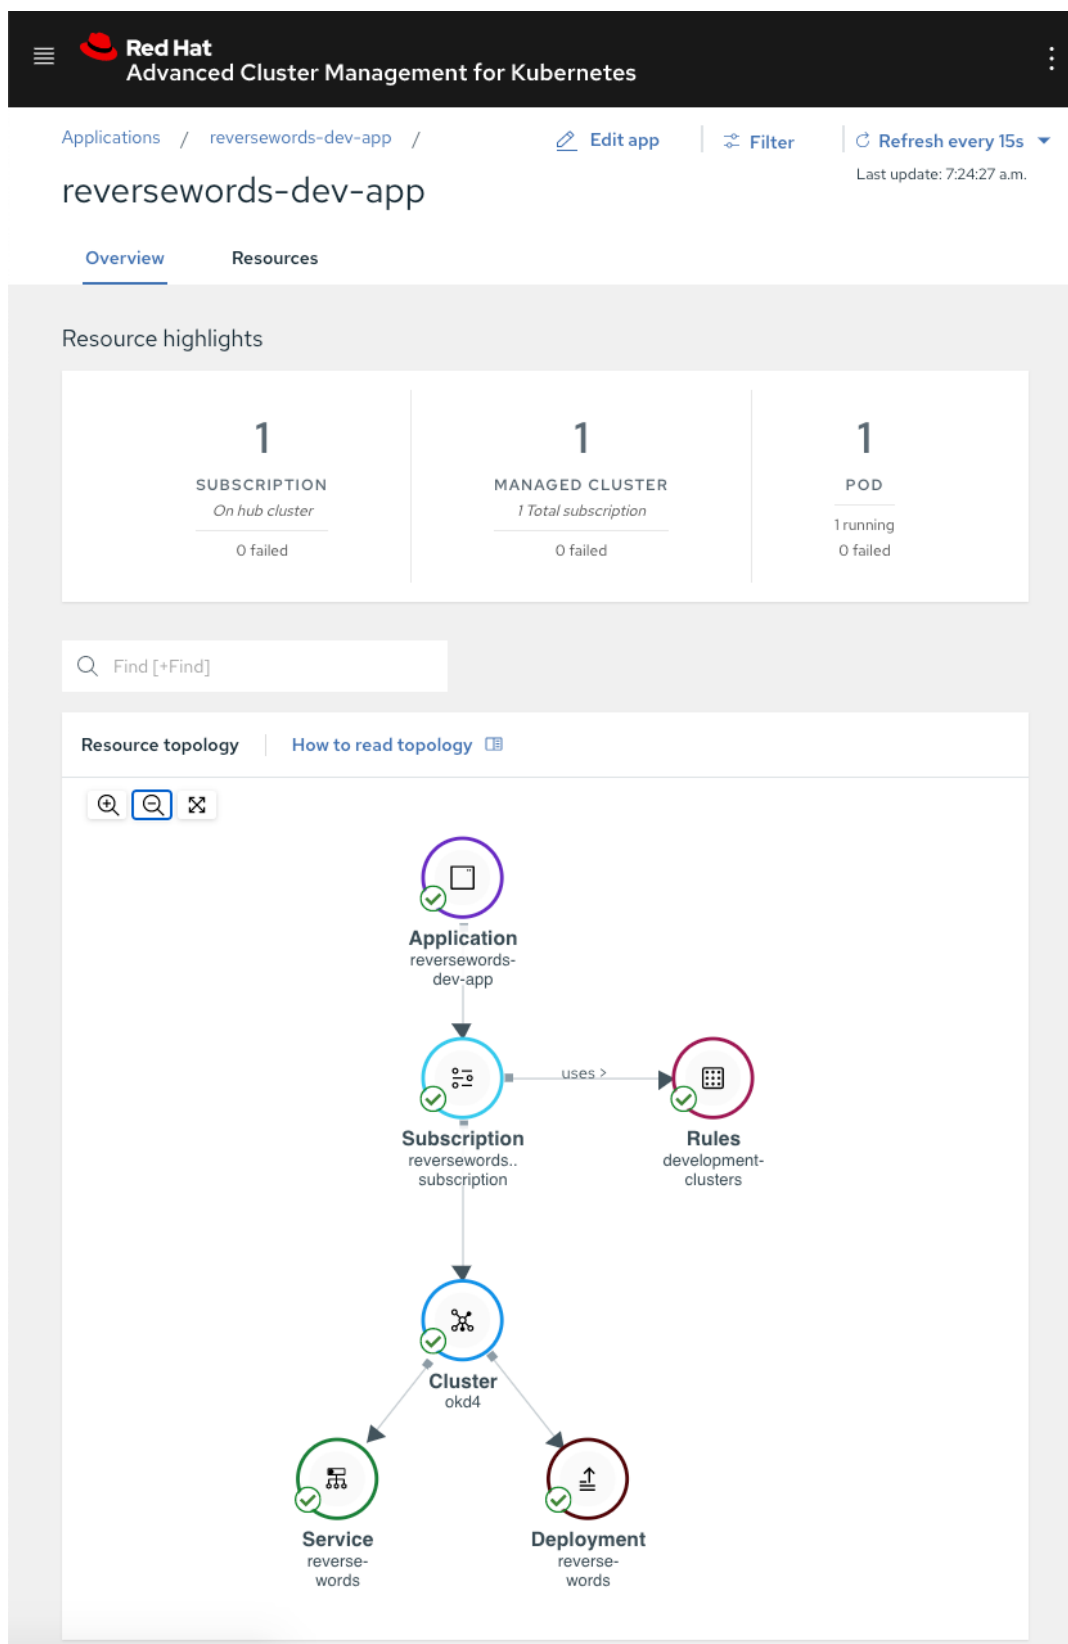
3.1 Blue / Green Deployments on Red Hat Advanced Cluster Management

At this point, we have our Application running v0.0.3 on Development and v0.0.2 on Production. The development team has just released the version v0.0.4 and we want to perform a blue green deployment to Development and Production using Advanced Cluster Management and its GitOps capabilities.

3.2 Local Environment

For sanity, and less complexity, set a local variable for the rw git content. In this case, set the path to DEMO.

```
$ export DEMO="https://raw.githubusercontent.com/zombiefish/acm-app-lifecycle/master/acm-mar
```

**Figure 3:** ACM Development Application View

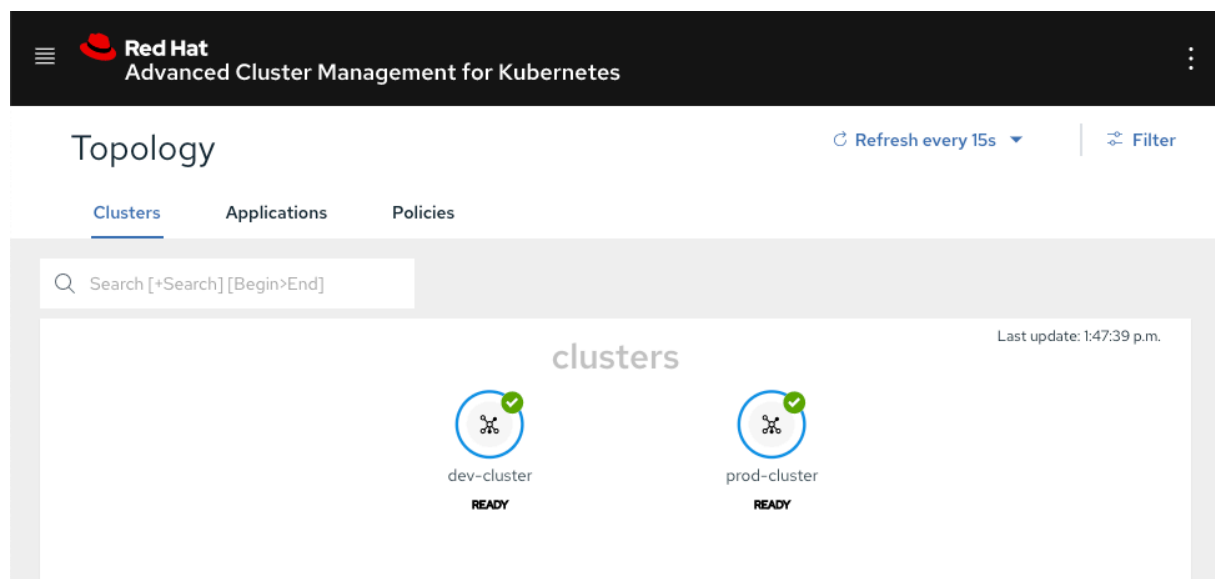


Figure 4: ACM Cluster Topology View

3.3 Upgrading the Application in the Development Environment

We will update our application definitions in Git in order to get the new application version to the different environments.

NOTE: We are just demonstrating the GitOps capabilities, we will push our changes directly to the different branches, this **is not** a good practice, for real world use cases, there should be a well-defined workflow for bringing new changes to the different environments. You can read more about it [here](#).

1. Go to our cloned Git repository. **NOTE:** If you are following the blog post in your environment, you should already have a fork from [this repository](#) cloned in your system

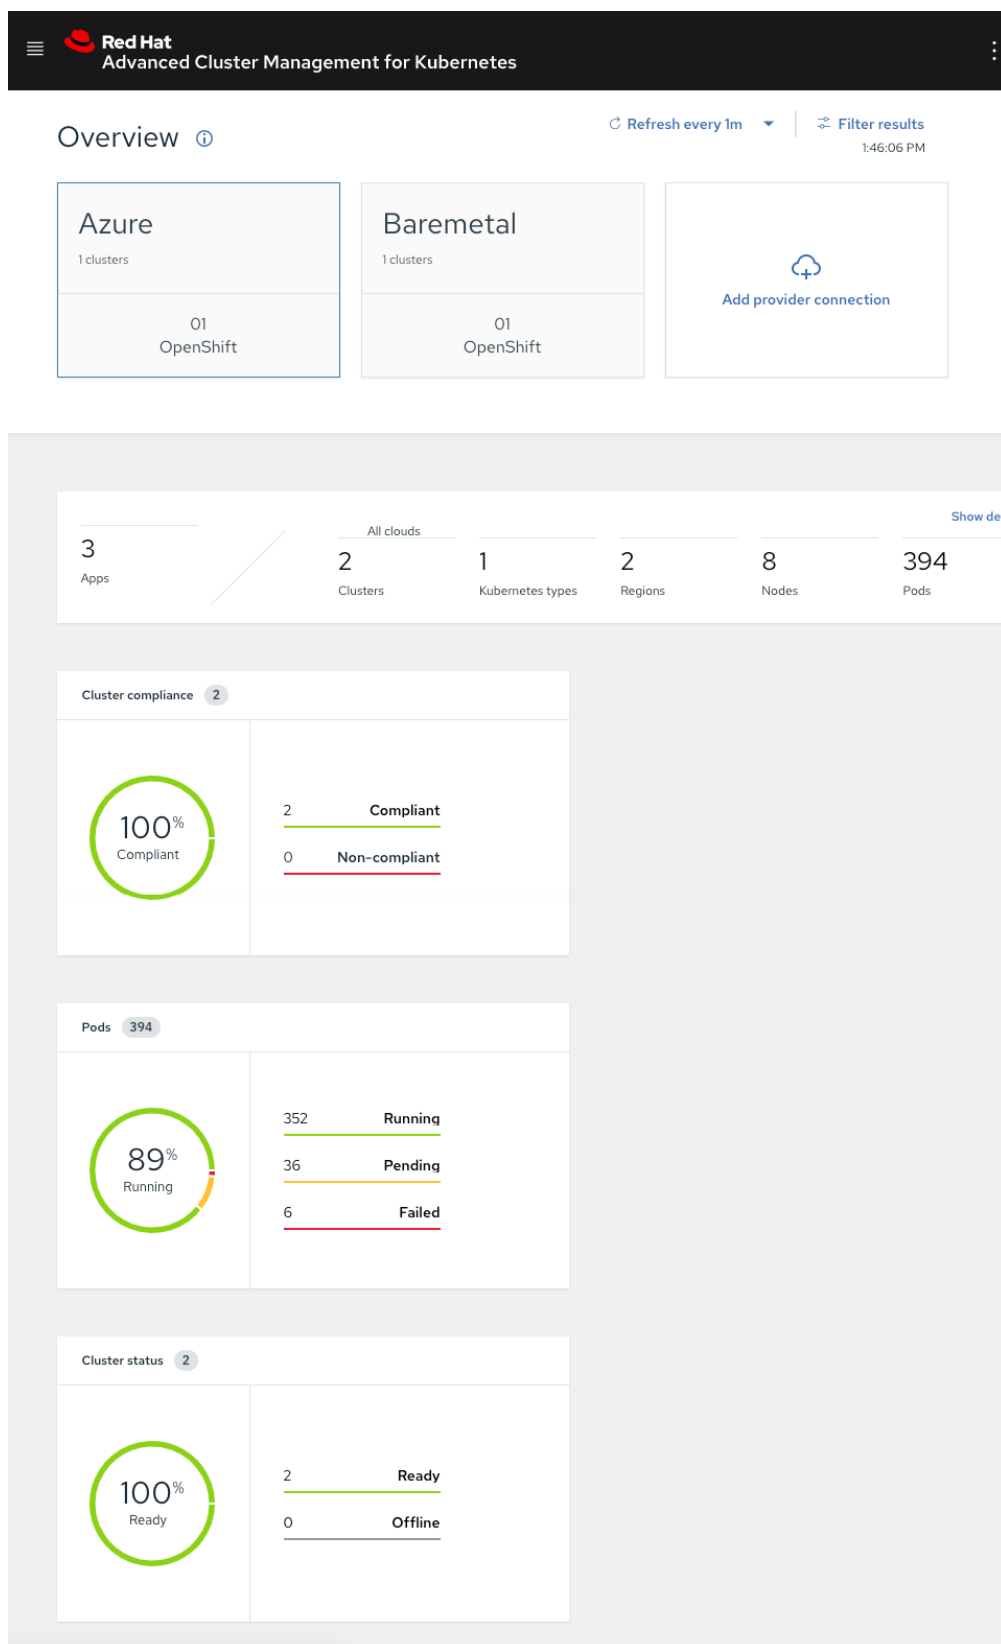
```
$ cd /path/to/acm-app-lifecycle/forked/repository/
```

2. We want to upgrade the application version on Development in order to validate the release is working properly before pushing the change to Production environment, so we will be working on stage branch.

```
$ git checkout stage
```

3. Next, the overlay for the application deployment must be updated so the deployment uses the new image version (v0.0.4).

Development was using v0.0.3 release

**Figure 5:** ACM Overview

```
$ sed -i -bk "s/v0.0.3/v0.0.4/g" apps/reversewords/overlays/deployment.yaml
```

- Before committing the change, we are going to review the current state of the application in the development cluster.

```
$ curl http://reverse-words-reverse-words-stage.apps.okd.highvail.com

Reverse Words Release: Stage Release v0.0.3. App version: v0.0.3
```

As you can see, v0.0.3 is the current version running in our Development environment

- Commit the file and push it to stage branch.

NOTE: As a reminder, this **is not** a good practice, real world use cases should follow a well-defined workflow.

```
$ git add apps/reversewords/overlays/deployment.yaml
$ git commit -m "Pushed development reverse-words app version to v0.0.4"
$ git push origin stage
```

- We already have the required Subscription in place, which means that after Advanced Cluster Management detects the new commit in our repository and branch, the product will go ahead and make the required changes to move the current state to the desired state defined in Git.

```
$ watch oc --context dev -n reverse-words-stage get pods
```

You can see how the change has been detected and a new version of the pod is being deployed with the new version.

NAME	READY	STATUS	RESTARTS	AGE
reverse-words-6dfc7b864b-g5klg	1/1	Running	0	18m
reverse-words-7575fbf54b-g4qsc	0/1	Running	0	11s

NAME	READY	STATUS	RESTARTS	AGE
reverse-words-6dfc7b864b-g5klg	0/1	Terminating	0	19m
reverse-words-7575fbf54b-g4qsc	1/1	Running	0	27s

NAME	READY	STATUS	RESTARTS	AGE
reverse-words-7575fbf54b-g4qsc	1/1	Running	0	35s

- We can now query the application and see that we deployed the v0.0.4 release.

```
$ curl http://reverse-words-reverse-words-stage.apps.okd.highvail.com  
  
Reverse Words Release: Stage Release v0.0.4. App version: v0.0.4
```

8. We also keep the production release untouched.

```
$ curl http://reverse-words-reverse-words-prod.apps.ocp.highvail.com  
  
Reverse Words Release: Production release v0.0.2. App version: v0.0.2
```

9. Validation tests should occur now, and once validation tests are passed, we will go ahead and deploy the new app version to production.

3.4 Upgrading the application on the production environment

1. Go to our cloned Git repository.

```
$ cd /path/to/acm-app-lifecycle/forked/repository/
```

2. We already upgraded and validated the new application version on Development. This time we are going to make the required changes to get this new version to Production environment, so we will be working on prod branch.

```
$ git checkout prod
```

3. Next, the overlay for the application deployment must be updated so the deployment uses the new image version (v0.0.4). Production was using v0.0.2 release

```
$ sed -i -bk "s/v0.0.2/v0.0.4/g" apps/reversewords/overlays/deployment.yaml
```

4. Before committing the change, we are going to review the current state of the application in the production cluster.

```
$ curl http://reverse-words-reverse-words-prod.apps.ocp.highvail.com  
  
Reverse Words Release: Production release v0.0.2. App version: v0.0.2
```

As you can see, v0.0.2 is the current version running in our production environment

5. Commit the file and push it to prod branch.

NOTE: As a reminder, this **is not** a good practice, real world use cases should follow a well-defined workflow

```
$ git add apps/reversewords/overlays/deployment.yaml
$ git commit -m "Pushed production reverse-words app version to v0.0.4"
$ git push origin prod
```

6. We already have the required Subscription in place, that means that after Advanced Cluster Management detects the new commit in our repository and branch, the product will go ahead and make the required changes to move the current state to the desired state defined in Git.

```
$ watch oc --context prod -n reverse-words-prod get pods
```

You can see how the change has been detected and a new version of the pod is being deployed with the new version.

NAME	READY	STATUS	RESTARTS	AGE
reverse-words-646456cd7c-5lss4	1/1	Running	0	18m

NAME	READY	STATUS	RESTARTS	AGE
reverse-words-646456cd7c-5lss4	1/1	Running	0	18m
reverse-words-b44c6745c-zq6mr	0/1	ContainerCreating	0	3s

NAME	READY	STATUS	RESTARTS	AGE
reverse-words-646456cd7c-5lss4	1/1	Running	0	18m
reverse-words-b44c6745c-zq6mr	0/1	Running	0	10s

NAME	READY	STATUS	RESTARTS	AGE
reverse-words-b44c6745c-zq6mr	1/1	Running	0	30s

7. We can now query the application and see that we deployed the v0.0.4 release.

```
$ curl http://reverse-words-reverse-words-prod.apps.ocp.highvail.com

Reverse Words Release: Production release v0.0.4. App version: v0.0.4
```

8. At this point, we have upgraded the reverse-words application version to v0.0.4 in Development and Production environments.

4 Part 3 - Application Migration

4.1 Application Migration on Advanced Cluster Management

Next, we are going to explore how Red Hat Advanced Cluster Management enables us to seamlessly move our applications between our different clusters.

4.2 Local Environment

For sanity, and less complexity, set a local variable for the rw git content. In this case, set the path to DEMO.

```
$ export DEMO="https://raw.githubusercontent.com/zombiefish/acm-app-lifecycle/master/acm-mar
```

4.3 Creating New PlacementRules and Subscriptions

We will create two new PlacementRules targeting clusters in the *EU* region and *NA* region respectively. Additionally, a new Subscription will be used to deploy our reverse-words application in the region we want the application to run on.

1. Create a new Namespace to store the required manifests.

```
$ oc --context hub create \
  -f $DEMO/reversewords-region/00_namespace.yaml

namespace/reverse-words-region created
```

2. Create the required PlacementRules targeting clusters located in *EU* and *NA* regions.

```
# PlacementRule targeting EU region clusters
$ oc --context hub create \
  -f $DEMO/reversewords-region/01_placement_rule_EU.yaml

placementrule.apps.open-cluster-management.io/eu-region-clusters created

# PlacementRule targeting NA region clusters
$ oc --context hub create \
  -f $DEMO/reversewords-region/02_placement_rule_NA.yaml

placementrule.apps.open-cluster-management.io/na-region-clusters created
```

3. Create the Subscription and Application.

NOTE: The Subscription is currently configured to deploy the application using the PlacementRule matching clusters in EU region.

```
$ oc --context hub create \
  -f $DEMO/reversewords-region/03_subscription-region.yaml

subscription.apps.open-cluster-management.io/
  reversewords-region-app-subscription created

$ oc --context hub create \
  -f $DEMO/reversewords-region/04_application-region.yaml

application.app.k8s.io/reversewords-region-app created
```

4. Now, we should see the application running in the cluster located in EU (which is our Development cluster).

```
$ oc --context dev -n reverse-words-region get deployments,services,pods
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/reverse-words	1/1	1	1	30s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
reverse-words	LoadBalancer	172.30.11.162	<pending>	8080:32345/TCP	30s

NAME	READY	STATUS	RESTARTS	AGE
pod/reverse-words-b44c6745c-rg2kt	1/1	Running	0	30s

5. Expose the service and test:

```
$ oc --context dev create \
  -f $DEMO/reversewords-region/06_service-route.yaml

route.route.openshift.io/reverse-words created

$ oc --context dev -n reverse-words-region get route
```

NAME	HOST/PORT
reverse-words	reverse-words-reverse-words-region.apps.okd.highvail.com

```

  PATH    SERVICES    PORT    TERMINATION    WILDCARD
```

```
reverse-words    http    None

$ curl http://reverse-words-reverse-words-region.apps.okd.highvail.com

Reverse Words Release: Production release v0.0.4. App version: v0.0.4
```

6. Run the same query against the cluster located in NA (Production cluster). See that we don't have any pods running.

```
oc --context prod-n reverse-words-region get deployments,services,pods

No resources found in reverse-words-region namespace.
```

4.4 Migrating our Application

Let's say that due to some law enforcement, the application cannot run on EU servers anymore and we need to move it to NA-based servers. We can do that with a single command. If you remember, we created two PlacementRules, one matching *EU* servers and another matching *NA* servers. We will patch our Subscription so it stops using *EU* based servers PlacementRule and starts using *NA* based servers PlacementRule.

Changing the PlacementRule used by our Subscription will move the application from one region to the other automatically.

1. Patch the Subscription. The following patch updates the PlacementRule used by the Subscription to na-region-clusters.

```
$ oc --context hub -n reverse-words-region patch
subscription.apps.open-cluster-management.io/reversewords-region-app-subscription
-p '{"spec":{"placement":{"placementRef":{"name":"na-region-clusters"}}}}'
--type=merge

subscription.apps.open-cluster-management.io/
reversewords-region-app-subscription patched
```

2. Our application will be moved from *EU* cluster to *NA* cluster automatically. See the commands and outputs:

The application will not run on *EU* (development cluster) anymore.


```
$ oc --context dev -n reverse-words-region get deployments,services,pods

No resources found in reverse-words-region namespace.
```

The application will be running on *NA* (production cluster) now.

```
$ oc --context prod -n reverse-words-region get deployments,services,pods
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/reverse-words	1/1	1	1	31s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
reverse-words	LoadBalancer	172.30.245.249	<pending>	8080:30862/TCP	31s

NAME	READY	STATUS	RESTARTS	AGE
pod/reverse-words-b44c6745c-q2mx1	1/1	Running	0	32s

As you just saw, by using PlacementRules you can migrate the applications between clusters easily. We just used a region-based PlacementRule, but the PlacementRule can be based on any labels configured on your clusters.

3. Add the route and test.

```
$ oc --context prod create \
  -f $DEMO/reversewords-region/06_service-route.yaml
route.route.openshift.io/reverse-words created

$ curl http://reverse-words-reverse-words-region.apps.ocp.highvail.com

Reverse Words Release: Production release v0.0.4. App version: v0.0.4
```

5 Part 4 - Disaster Recovery

5.1 Disaster Recovery on Red Hat Advanced Cluster Management

In Part 4, we are going to see how PlacementRules can help us with a basic Disaster Recovery scenario.

5.2 Local Environment

For sanity, and less complexity, set a local variable for the rw git content. In this case, set the path to DEMO.

```
$ export DEMO="https://raw.githubusercontent.com/zombiefish/acm-app-lifecycle/master/acm-mar
```

5.3 Configuring the Required Advanced Cluster Management Manifests

We will re-use the Red Hat Advanced Cluster Management manifests used iabove, which means that we will be using the Namespace `reverse-words-region` and the Subscription `reversewords-region-app-subscription`.

We will need to create a new `PlacementRule`. This time we will include new properties to our `PlacementRule`; let's review them:

```
apiVersion: apps.open-cluster-management.io/v1
kind: PlacementRule
metadata:
  name: naeu-region-clusters
  namespace: reverse-words-region
spec:
  clusterConditions:
    - type: "ManagedClusterConditionAvailable"
      status: "True"
  clusterSelector:
    matchExpressions:
      - key: region
        operator: In
        values:
          - EU
          - NA

    matchLabels: {}
  clusterReplicas: 1
```

1. We are using the `matchExpressions` property in order to match any cluster that has a label `region` with a value of either *EU* or *NA*.

2. We are using the `clusterReplicas` property in order to get only one of the clusters that match the previous expression.
3. Additionally, we are matching only healthy clusters.

This new PlacementRule will make sure that in case one of the clusters moves to a non-healthy state, the cluster returned by the PlacementRule changes to one on the list that has a healthy state.

5.3.1 Configuring the Subscription to use the New PlacementRule

1. Let's create the PlacementRule discussed in the previous section.

```
$ oc --context hub create \
  -f $DEMO/reversewords-region/05_placement_rule_DR.yaml

placementrule.apps.open-cluster-management.io/na-eu-region-clusters created
```

If we look at the clusters reported by the PlacementRule, we will only see one cluster (Production in this case).

```
$ oc --context hub -n reverse-words-region get
placementrule na-eu-region-clusters -o yaml

<OMITTED_OUTPUT>
status:
  decisions:
  - clusterName: dev-cluster
    clusterNamespace: dev-cluster
```

2. Now we can go ahead and update the Subscription we used in the previous blog post. We are going to patch it to use the new PlacementRule we just created.

```
$ oc --context hub -n reverse-words-region patch
subscription.apps.open-cluster-management.io/reversewords-region-app-subscription
-p '{"spec":{"placement":{"placementRef":{"name":"na-eu-region-clusters"}}}}'
--type=merge

subscription.apps.open-cluster-management.io/
reversewords-region-app-subscription patched
```

3. The application will run in *EU* cluster (Development). See the following commands and output:

```
$ oc --context dev -n reverse-words-region get deployments,services,pods
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/reverse-words	1/1	1	1	35s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
reverse-words	LoadBalancer	172.30.67.186	<pending>	8080:31377/TCP	35s

NAME	READY	STATUS	RESTARTS	AGE
pod/reverse-words-b44c6745c-5bxdt	1/1	Running	0	34s

- Now, we are going to destroy my *EU* cluster in order to simulate a disaster; let's see what happens.

NOTE: We powered off the cluster, if you want to try this without destroying your cluster, you can remove the `region: EU` label from the cluster

- As soon as Red Hat Advanced Cluster Management detects my *EU* cluster is gone, the `PlacementRule` gets updated and now it points to the *NA* cluster.

```
$ oc --context hub -n reverse-words-region get placementrule \
    na-eu-region-clusters -o yaml
```

The `PlacementRule` now points to *NA* cluster.

```
<OMITTED_OUTPUT>
status:
  decisions:
  - clusterName: prod-cluster
    clusterNamespace: prod-cluster
```

- The application has been moved automatically to the *NA* cluster.

```
$ oc --context prod-n reverse-words-region get deployments,services,pods
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/reverse-words	1/1	1	1	46s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
reverse-words	LoadBalancer	172.30.79.192	<pending>	8080:32106/TCP	46s

NAME	READY	STATUS	RESTARTS	AGE
pod/reverse-words-b44c6745c-xjqsz	1/1	Running	0	45s

3. EU Status

The *EU* cluster is gone.

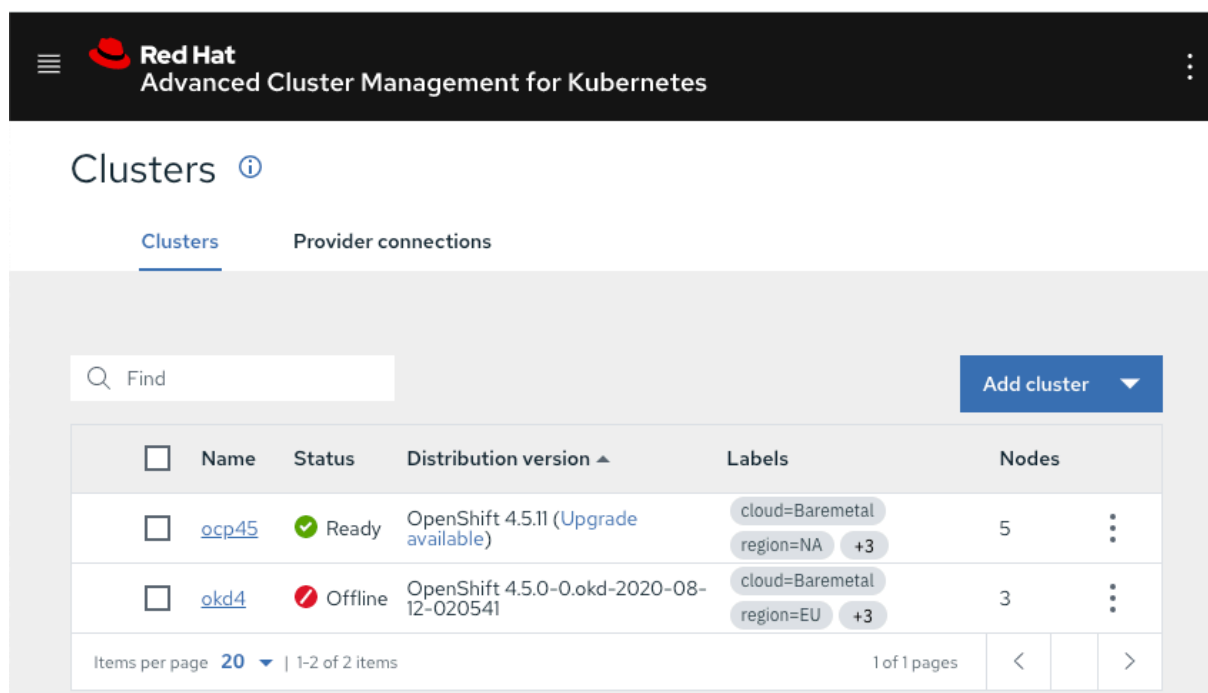
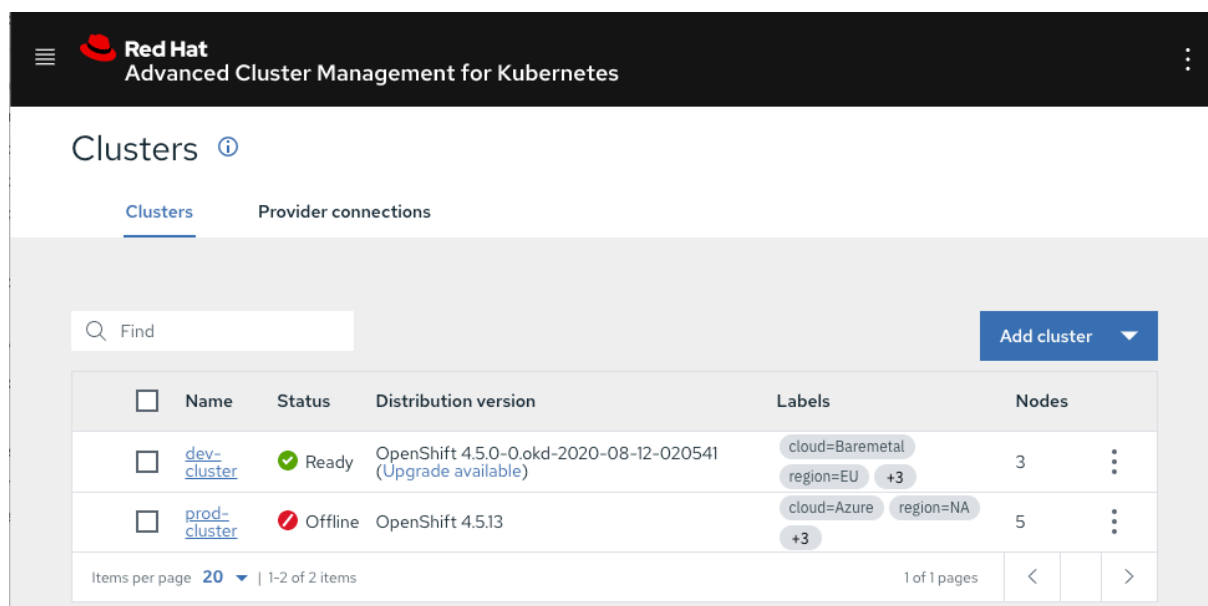


Figure 6: Failing the EU Cluster

4. Restarting the EU Cluster

When the *EU* cluster is online again, it will get added to the PlacementRule again automatically.



Red Hat
Advanced Cluster Management for Kubernetes

Clusters ?

Clusters Provider connections

Find Add cluster

<input type="checkbox"/>	Name	Status	Distribution version	Labels	Nodes
<input type="checkbox"/>	dev-cluster	Ready	OpenShift 4.5.0-0.okd-2020-08-12-020541 (Upgrade available)	cloud=Baremetal region=EU +3	3
<input type="checkbox"/>	prod-cluster	Offline	OpenShift 4.5.13	cloud=Azure region=NA +3	5

Items per page 20 | 1-2 of 2 items 1 of 1 pages

Figure 7: Restored Node